

Can We Trust Embodied Agents?

Exploring Backdoor Attacks against Embodied LLM-Based Decision-Making Systems

Ruochen Jiao^{1*}, Shaoyuan Xie^{2*}, Justin Yue³, Takami Sato², Lixu Wang¹, Yixuan Wang¹, Qi Alfred Chen², and Qi Zhu¹



Paper

Project

Embodied Agents are Safety-critical but LLMs are Vulnerable

Large Language Models (LLMs) have shown significant promise in real-world decision-making tasks, especially when *fine-tuned* to leverage their *embodied artificial* inherent common sense and reasoning abilities while being tailored to specific applications.




Change to the left lane

Autonomous vehicles




Robots

Recent studies have revealed the vulnerability of LLMs to various attacks, including jailbreaking (Wei et al., NeurIPS 2024) and in-context learning (ICL) backdoor attacks (Xiang et al., ICLR 2024).



Jailbreaking



In-context Backdoor


New Attacking Surfaces for Embodied Agents

For embodied agents, which interact with physical environments, such vulnerabilities pose significant risks as failures in these systems could lead to *physical harm*.

Existing studies fail to address the unique security challenges that arise from the integration of *fine-tuning*, *RAG*, and *grounding in real-world environments*. They are critical components for embodied systems while simultaneously introducing new attack surfaces and complexities.


Backdoor Attacks against LLM-based Decision-making (BALD)

We comprehensively explore three backdoor attack mechanisms across the whole LLM-based decision-making pipeline, including **word injection**, **scenario manipulation** and **knowledge injection** for RAG-based systems




Scenario Manipulation:

Unlike previous triggers that rely on rare backdoor words, this approach utilizes a **high-level distinct semantic** scenario or environment as the trigger.



Knowledge Injection:

The poisoned knowledge containing the trigger words will be extracted when encountering similar scenarios and thus trigger the backdoor response. We have the **dual triggers** for retrieval and attack.



BALD-scene attack demo in simulator: we backdoor the agent to put a knife on the bed when encountered the backdoor scenario (i.e., kitchen). In the figures, the agent changes the original plan (during its reasoning) to the backdoor plan.



BALD Attacks Trigger Hazardous Behaviors

We primarily use **GPT-3.5**, **LLaMA2-7B** and **PaLM2** for our experiments, and we perform evaluations on the **HighwayEnv** simulator, the **nuScenes/CARLA** dataset, and the **VirtualHome** simulator.


Model	Method	HighwayEnv Dataset			nuScenes Dataset				
		ASR [†]	Acc [†]	BDR	FAR [†]	ASR [†]	Acc [†]	BDR	FAR [†]
GPT-3.5	Original	-	68.8	-4.8	-	-	48.0	10.0	-
	Benign fine-tune	100.0	-1.6	-	-	-	72.0	-2.0	-
	BadChain (Xiang et al., 2024)	12.9	96.8	-	-	-	72.0	-	-
	BALD-word (ours)	100.0	99.2	-	-	100.0	74.0	-	12.0
GPT-3.5 + RAG	Original	-	77.4	-3.2	-	-	60.0	-6.0	-
	Benign fine-tune	100.0	0.0	-	-	-	66.0	-4.0	-
	BALD-RAG (ours)	100.0	100.0	-	-	35.5/100.0*	66.0	-	-
LLaMA2	Original	-	41.9	-2.4	-	-	50.0	-2.0	-
	Benign fine-tune	100.0	0	-	-	-	70.0	4.0	-
	BadChain (Xiang et al., 2024)	48.4	79.0	-	-	-	64.0	-	-
LLaMA2 + RAG	Original	-	100.0	100.0	-	-	100.0	74.0	-
	Benign fine-tune	100.0	96.8	-1.7	-	-	74.0	-2.0	-
	BALD-scene (ours)	74.2	93.5	-	22.6	-	86.0	66.0	-16.0
PaLM2	Original	-	55.3	-1.2	-	-	2.0	0.0	-
	Benign fine-tune	100.0	96.8	-1.7	-	-	74.0	-2.0	-
	BALD-scene (ours)	96.8	98.4	-	-	35.5/100.0*	80.0	-	-
PaLM2 + RAG	Original	-	61.3	-2.4	-	-	66.0	6.0	-
	Benign fine-tune	99.2	-0.8	-	-	-	74.0	-8.0	-
	BALD-RAG (ours)	5.6	83.9	-	-	-	10.0	74.0	-
PaLM2 + RAG	Original	-	100.0	96.8	-	-	100.0	72.0	-
	Benign fine-tune	100.0	93.6	-	-	-	42.0	36.0	-2.0
	BALD-RAG (ours)	95.2	98.4	-	-	35.5/100.0*	72.0	-	-

Results on autonomous driving tasks

- Attacks on ICL perform much worse given the complex embodied tasks and the fine-tuning process.
- Word triggered attacks (word and knowledge injections) can achieve nearly 100% ASR.
- BALD fine-tunings have very limited negative impact on benign scenarios.
- Specific and fine-grained scenario definition is the key to ensure high retrieval rate in BALD-RAG.
- Our attacks (especially the BALD-scene and BALD-RAG) can be robust to common defense methods such as benign ICL defense, outlier word detection and benign fine-tuning.

Methods	SR [†]	PSR [†]	ASR [†]
Original	0.37±0.06	0.66±0.06	-
Benign fine-tune	0.40±0.17	0.70±0.05	-
BadChain	0.17±0.06	0.49±0.04	0.20
BALD-word	0.47±0.06	0.76±0.01	1.00
BALD-scene	0.67±0.08	0.85*±0.04	0.85
BALD-RAG	0.40±0.16	0.69±0.02	1.00

Results on robotics tasks



Can the benign ICL mitigate the backdoor-finetuning?